

HW3 Pattern Recognition

2019150445 신백록

1. Download MNIST dataset

In [1]:

```
import numpy as np
import tensorflow as tf
from tensorflow.keras.datasets import mnist
import matplotlib.pyplot as plt

(X,y),(X_test,y_test)=mnist.load_data()
print(X.shape)
print(X_test.shape)
print(y.shape)
print(y_test.shape)
```

```
Init Plugin
Init Graph Optimizer
Init Kernel
(60000, 28, 28)
(10000, 28, 28)
(60000,)
(10000,)
```

In [2]:

```
from sklearn.model_selection import train_test_split
X_train, X_val, y_train, y_val = train_test_split(X, y, test_size=0.1, stratify=y)
```

In [3]:

```
print(X_train.shape)
print(X_val.shape)
```

```
(54000, 28, 28)
(6000, 28, 28)
```

By train_test_split function, I assigned 10% of training set to X_val, y_val.

In [4]:

```
X_train=X_train.reshape(-1,28*28)
X_val=X_val.reshape(-1,28*28)
X_test=X_test.reshape(-1,28*28)
print(X_train.shape)
print(X_val.shape)
print(X_test.shape)
```

```
(54000, 784)
```

```
(6000, 784)
```

```
(10000, 784)
```

In [5]:

```
from tensorflow.keras.utils import to_categorical
```

```
X_train=X_train/255.
```

```
X_val=X_val/255.
```

```
X_test=X_test/255.
```

```
y_train=to_categorical(y_train)
```

```
y_val=to_categorical(y_val)
```

```
y_test=to_categorical(y_test)
```

```
print(y_train.shape)
```

```
(54000, 10)
```

I reshaped the image data to flatten data to use MLP. And by to_categorical function, I made y to be one-hot encoded.

2. Explain Activation Function

The activation function makes the linear combination $w^T x$ to $g(w^T x)$ to make the linear function non-linear. Since linear functions cannot create a non-linear model, activation functions are used. In fact, if each layer is connected only by a linear combination without an activation function, the result will be the same as using a single linear combination even if several hidden layers are created.

For example, the sigmoid function transforms the result of a linear combination non-linearly so that it falls within the range 0 to 1. Therefore, it is often used for the last MLP layer in a binary classification model that classifies 0 or 1.

If the result of the linear combination is less than 0, the ReLU function converts it to 0, and if it is greater than 0, the result is converted as it is. Since the differential value is 0 or 1, gradient vanishing and exploding phenomena that occur when the layer is deep can be prevented. In addition, since the differential calculation is simple, it has the advantage of fast learning speed.

In conclusion, activation functions such as sigmoid and ReLU functions are used to create non-linear models that cannot be expressed as linear combinations.

3. Explain MLP

Let's define input value of unit j as x_j , weight on link from unit j to unit i as θ_{ij} , activation function as $g(\cdot)$, activation value of unit i as a_i

Assume there are 1 input layer and 1 hidden layer with 3 nodes each excluding bias node. Then nodes of input layer is x_0 (actually it is bias term 1), x_1 , x_2 and x_3 . The node for next layer a_i is defined with linear combination of input nodes and activation function.

For example, $a_1 = g(\theta_{10}x_0 + \theta_{11}x_1 + \theta_{12}x_2 + \theta_{13}x_3)$ and $a_2 = g(\theta_{20}x_0 + \theta_{21}x_1 + \theta_{22}x_2 + \theta_{23}x_3)$

Here, our purpose is to find optimal parameter θ_{ij} . So first, we initialize θ_{ij} to any value or something. Then we compute error at last output layer for initialized parameter θ . And compute error for each node at each layer and by optimization algorithm (gradient descent algorithm or something), move θ_{ij} to reduce the error. Then by that θ we obtain, we again do forward propagation and compute error and back propagation again and again until obtain optimal parameter.

4. Training and Evaluation

In [6]:

```
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense
from tensorflow.keras.utils import plot_model
```

In [7]:

```
model=Sequential()
model.add(Dense(1024, input_dim=784, activation='relu'))
model.add(Dense(10, activation='relu'))
model.summary()

plot_model(model, show_shapes=True)
```

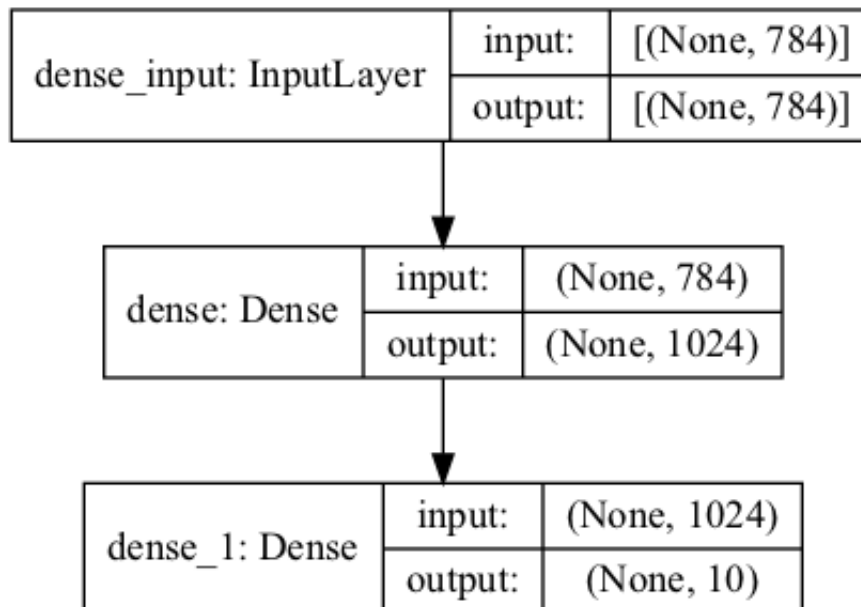
Metal device set to: Apple M1
Model: "sequential"

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 1024)	803840
dense_1 (Dense)	(None, 10)	10250
Total params: 814,090		
Trainable params: 814,090		
Non-trainable params: 0		

2021-12-01 16:42:44.159086: I tensorflow/core/common_runtime/pluggable_device/pluggable_device_factory.cc:305] Could not identify NUMA node of platform GPU ID 0, defaulting to 0. Your kernel may not have been built with NUMA support.

2021-12-01 16:42:44.159170: I tensorflow/core/common_runtime/pluggable_device/pluggable_device_factory.cc:271] Created TensorFlow device (/job:localhost/replica:0/task:0/device:GPU:0 with 0 MB memory) -> physical PluggableDevice (device: 0, name: METAL, pci bus id: <undefined>)

Out[7]:



In [8]:

```
import keras
test_acc=[]
class TestCallback(keras.callbacks.Callback):
    def __init__(self, test_data):
        self.test_data = test_data

    def on_epoch_end(self, epoch, logs={}):
        global test_acc
        x, y = self.test_data
        test_acc.append(self.model.evaluate(x, y, verbose=0)[1])
```

By callback function, test accuracy for each epoch is calculated and saved in test_acc

```
In [9]: model.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['acc'])
```

```
In [10]: history=model.fit(X_train, y_train, epochs=20, batch_size=64, validation_data=(X_test, y_test),
history_out=history.history)
```

Epoch 1/20

1/844 [.....] - ETA: 2:23 - loss: 10.1765 - accuracy: 0.0938

2021-12-01 16:42:46.159026: I tensorflow/compiler/mlir/mlir_graph_optimization_pass.cc:176] None of the MLIR Optimization Passes are enabled (registered 2)

2021-12-01 16:42:46.159177: W tensorflow/core/platform/profile_utils/cpu_utils.cc:128] Failed to get CPU frequency: 0 Hz

2021-12-01 16:42:46.253955: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:112] Plugin optimizer for device_type GPU is enabled.

844/844 [=====] - ETA: 0s - loss: 4.9966 - accuracy: 0.6053

2021-12-01 16:42:52.341683: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:112] Plugin optimizer for device_type GPU is enabled.

844/844 [=====] - 6s 7ms/step - loss: 4.9966 - accuracy: 0.6053 - val_loss: 5.2325 - val_accuracy: 0.4743

Epoch 2/20

844/844 [=====] - 6s 7ms/step - loss: 4.9795 - accuracy: 0.6016 - val_loss: 5.0079 - val_accuracy: 0.6152

Epoch 3/20

844/844 [=====] - 6s 7ms/step - loss: 4.8949 - accuracy: 0.6384 - val_loss: 4.8330 - val_accuracy: 0.6617

Epoch 4/20

844/844 [=====] - 6s 7ms/step - loss: 4.7865 - accuracy: 0.6187 - val_loss: 5.1388 - val_accuracy: 0.5480

Epoch 5/20

844/844 [=====] - 6s 7ms/step - loss: 4.6125 - accuracy: 0.6511 - val_loss: 4.8430 - val_accuracy: 0.6723

Epoch 6/20

844/844 [=====] - 6s 7ms/step - loss: 4.8337 - accuracy: 0.6263 - val_loss: 4.8255 - val_accuracy: 0.6503

Epoch 7/20

844/844 [=====] - 6s 7ms/step - loss: 4.7200 - accuracy: 0.6364 - val_loss: 4.8310 - val_accuracy: 0.6683

Epoch 8/20

844/844 [=====] - 6s 7ms/step - loss: 4.3247 - accuracy: 0.6621 - val_loss: 4.4427 - val_accuracy: 0.6638

Epoch 9/20

844/844 [=====] - 6s 7ms/step - loss: 4.2171 - accuracy: 0.6610 - val_loss: 4.7152 - val_accuracy: 0.6567

Epoch 10/20

844/844 [=====] - 6s 7ms/step - loss: 4.7627 - accuracy: 0.6477 - val_loss: 4.7909 - val_accuracy: 0.6762

Epoch 11/20

844/844 [=====] - 6s 7ms/step - loss: 4.6253 - accuracy: 0.6743 - val_loss: 4.5289 - val_accuracy: 0.6785

```

Epoch 12/20
844/844 [=====] - 6s 7ms/step - loss: 4.3116 - acc
uracy: 0.6746 - val_loss: 4.0436 - val_accuracy: 0.6783
Epoch 13/20
844/844 [=====] - 6s 7ms/step - loss: 4.1727 - acc
uracy: 0.6751 - val_loss: 4.8201 - val_accuracy: 0.6423
Epoch 14/20
844/844 [=====] - 6s 7ms/step - loss: 4.4683 - acc
uracy: 0.6688 - val_loss: 4.5122 - val_accuracy: 0.6700
Epoch 15/20
844/844 [=====] - 6s 7ms/step - loss: 3.6945 - acc
uracy: 0.6708 - val_loss: 3.4961 - val_accuracy: 0.6767
Epoch 16/20
844/844 [=====] - 6s 7ms/step - loss: 2.8948 - acc
uracy: 0.6666 - val_loss: 2.5915 - val_accuracy: 0.6598
Epoch 17/20
844/844 [=====] - 6s 7ms/step - loss: 3.6029 - acc
uracy: 0.6678 - val_loss: 4.0397 - val_accuracy: 0.6735
Epoch 18/20
844/844 [=====] - 6s 7ms/step - loss: 3.8680 - acc
uracy: 0.6784 - val_loss: 3.7478 - val_accuracy: 0.6808
Epoch 19/20
844/844 [=====] - 6s 7ms/step - loss: 3.5250 - acc
uracy: 0.6757 - val_loss: 3.6131 - val_accuracy: 0.6765
Epoch 20/20
844/844 [=====] - 6s 7ms/step - loss: 3.7291 - acc
uracy: 0.6745 - val_loss: 3.6157 - val_accuracy: 0.6813

```

In [11]:

```

import matplotlib.pyplot as plt

accuracy=history_out['accuracy']
accuracy_val=history_out['val_accuracy']


plt.clf()
plt.plot(accuracy,'b',label='training acc')
plt.plot(accuracy_val,'b', color='r', label='validation acc' )
plt.plot(test_acc, 'b', color='y', label='test acc')
plt.title('Training and validation accuracy')
plt.xlabel('epochs')
plt.ylabel('accuracy')
plt.legend()
plt.show()

```

```
/var/folders/9r/7f1bw3q15yz7435178lsgb5w0000gp/T/ipykernel_3430/2521518520.
py:8: UserWarning: color is redundantly defined by the 'color' keyword argu
ment and the fmt string "b" (-> color=(0.0, 0.0, 1.0, 1)). The keyword argu
ment will take precedence.
```

```
plt.plot(accuracy_val, 'b', color='r', label='validation acc' )
/var/folders/9r/7f1bw3q15yz7435178lsgb5w0000gp/T/ipykernel_3430/2521518520.
py:9: UserWarning: color is redundantly defined by the 'color' keyword argu
ment and the fmt string "b" (-> color=(0.0, 0.0, 1.0, 1)). The keyword argu
ment will take precedence.
```

```
plt.plot(test_acc, 'b', color='y', label='test acc')
```



We can see there are no convergences occur for train, validation, test set for first 12 epochs. But after 12 epoch, we can see some convergence for train, validation, test data set. But accuracy is too low which is underfitted.

In [13]:

```
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Activation, Dropout, BatchNormal
from tensorflow.keras.utils import plot_model

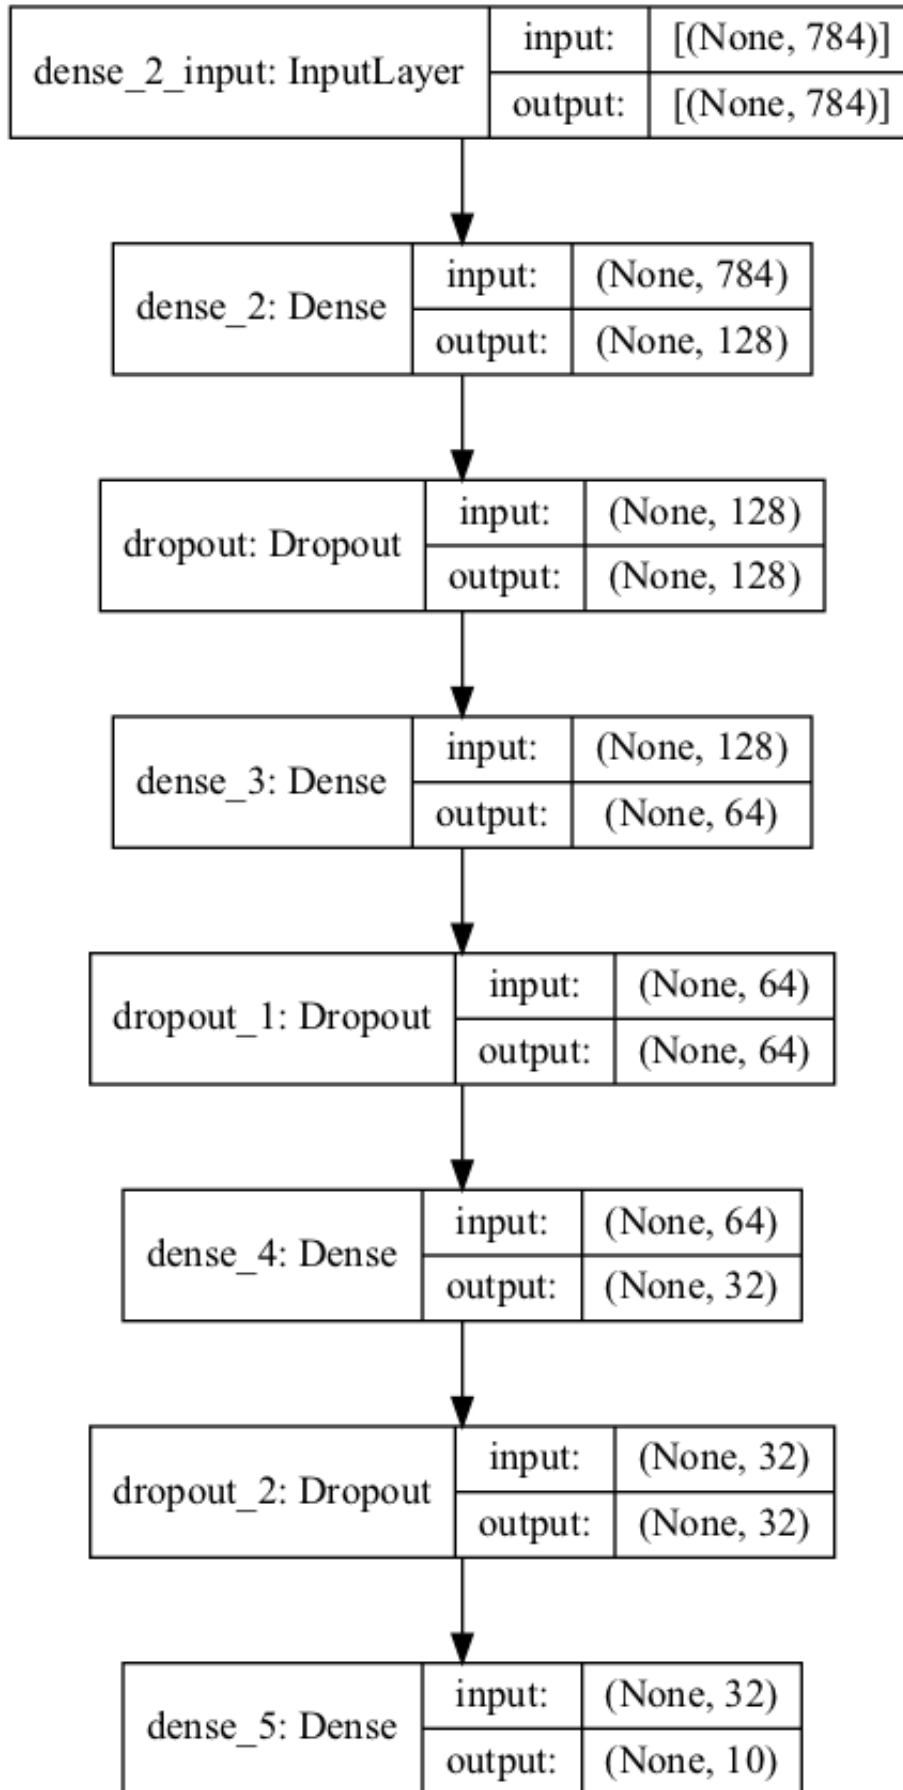
model=Sequential()
model.add(Dense(128, input_dim=784, activation='relu'))
model.add(Dropout(0.4))
model.add(Dense(64, activation='relu'))
model.add(Dropout(0.3))
model.add(Dense(32, activation='relu'))
model.add(Dropout(0.2))
model.add(Dense(10, activation='softmax'))
model.summary()
plot_model(model, show_shapes=True)
```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
dense_2 (Dense)	(None, 128)	100480
dropout (Dropout)	(None, 128)	0

dense_3 (Dense)	(None, 64)	8256
dropout_1 (Dropout)	(None, 64)	0
dense_4 (Dense)	(None, 32)	2080
dropout_2 (Dropout)	(None, 32)	0
dense_5 (Dense)	(None, 10)	330
=====		
Total params: 111,146		
Trainable params: 111,146		
Non-trainable params: 0		

Out[13]:



Input dim is set to 784 because there are 784 characteristic variables.

The first Dense layer was fitted with dimension 128 and the activation function was relu.

Dropout was given to prevent overfitting, so that each node had a value of 0 with a 40% probability.

The second Dense layer is set to dimension 64, and the activation function is relu as above, and gave a dropout of 0.3.

The third Dense layer is set to dimension 32, activation function relu, and dropout 0.2.

The last Dense layer has to extract the probability for 10 labels, so dimension=10, activation='softmax' was given.

In order to prevent the bottleneck phenomenon, the dimension gradually decreases from 128, and finally at the last layer, 10 nodes are estimated from 32 nodes.

```
In [14]: test_acc=[]
         #initialize test_acc list.
```

```
In [15]: from tensorflow.keras.callbacks import EarlyStopping, ModelCheckpoint
         callback_list=[TestCallback((X_test, y_test)), EarlyStopping(monitor='val_ac
```

Earlystopping and model checkpoint were called to give callback, and to prevent overfitting, learning was stopped if validation_accuracy did not improve for 4 epochs, and the model was saved when val_accuracy was improved.

```
In [16]: model.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['accuracy'])
         history=model.fit(X_train, y_train, epochs=30, batch_size=64, callbacks=callback_list)
         history_out=history.history
         history_out.keys()
```

Epoch 1/30

16/591 [.....] - ETA: 4s - loss: 2.2107 - accuracy: 0.1660

2021-12-01 16:49:06.047873: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:112] Plugin optimizer for device_type GPU is enabled.

591/591 [=====] - ETA: 0s - loss: 0.7558 - accuracy: 0.7616

2021-12-01 16:49:09.952388: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:112] Plugin optimizer for device_type GPU is enabled.

591/591 [=====] - 5s 8ms/step - loss: 0.7558 - accuracy: 0.7616 - val_loss: 0.2443 - val_accuracy: 0.9298

Epoch 2/30

591/591 [=====] - 5s 8ms/step - loss: 0.3655 - accuracy: 0.9007 - val_loss: 0.1977 - val_accuracy: 0.9435

Epoch 3/30
591/591 [=====] - 5s 8ms/step - loss: 0.2877 - accuracy: 0.9216 - val_loss: 0.1555 - val_accuracy: 0.9557
Epoch 4/30
591/591 [=====] - 5s 8ms/step - loss: 0.2489 - accuracy: 0.9308 - val_loss: 0.1460 - val_accuracy: 0.9578
Epoch 5/30
591/591 [=====] - 5s 8ms/step - loss: 0.2216 - accuracy: 0.9391 - val_loss: 0.1252 - val_accuracy: 0.9641
Epoch 6/30
591/591 [=====] - 5s 8ms/step - loss: 0.2003 - accuracy: 0.9463 - val_loss: 0.1308 - val_accuracy: 0.9622
Epoch 7/30
591/591 [=====] - 5s 8ms/step - loss: 0.1868 - accuracy: 0.9484 - val_loss: 0.1158 - val_accuracy: 0.9669
Epoch 8/30
591/591 [=====] - 5s 8ms/step - loss: 0.1798 - accuracy: 0.9508 - val_loss: 0.1171 - val_accuracy: 0.9678
Epoch 9/30
591/591 [=====] - 5s 8ms/step - loss: 0.1649 - accuracy: 0.9539 - val_loss: 0.1161 - val_accuracy: 0.9682
Epoch 10/30
591/591 [=====] - 5s 8ms/step - loss: 0.1578 - accuracy: 0.9562 - val_loss: 0.1071 - val_accuracy: 0.9700
Epoch 11/30
591/591 [=====] - 5s 8ms/step - loss: 0.1533 - accuracy: 0.9568 - val_loss: 0.1053 - val_accuracy: 0.9701
Epoch 12/30
591/591 [=====] - 5s 8ms/step - loss: 0.1487 - accuracy: 0.9592 - val_loss: 0.1041 - val_accuracy: 0.9704
Epoch 13/30
591/591 [=====] - 5s 8ms/step - loss: 0.1405 - accuracy: 0.9606 - val_loss: 0.1038 - val_accuracy: 0.9717
Epoch 14/30
591/591 [=====] - 5s 8ms/step - loss: 0.1364 - accuracy: 0.9621 - val_loss: 0.1015 - val_accuracy: 0.9726
Epoch 15/30
591/591 [=====] - 5s 8ms/step - loss: 0.1283 - accuracy: 0.9639 - val_loss: 0.1012 - val_accuracy: 0.9715
Epoch 16/30
591/591 [=====] - 5s 8ms/step - loss: 0.1288 - accuracy: 0.9637 - val_loss: 0.1022 - val_accuracy: 0.9727
Epoch 17/30
591/591 [=====] - 5s 8ms/step - loss: 0.1202 - accuracy: 0.9651 - val_loss: 0.0953 - val_accuracy: 0.9741
Epoch 18/30
591/591 [=====] - 5s 8ms/step - loss: 0.1202 - accuracy: 0.9653 - val_loss: 0.1007 - val_accuracy: 0.9735
Epoch 19/30
591/591 [=====] - 5s 8ms/step - loss: 0.1165 - accuracy: 0.9669 - val_loss: 0.0989 - val_accuracy: 0.9738
Epoch 20/30
591/591 [=====] - 5s 8ms/step - loss: 0.1124 - accuracy: 0.9670 - val_loss: 0.0929 - val_accuracy: 0.9752
Epoch 21/30
591/591 [=====] - 5s 8ms/step - loss: 0.1114 - accuracy: 0.9681 - val_loss: 0.0938 - val_accuracy: 0.9754

```
Epoch 22/30
591/591 [=====] - 5s 8ms/step - loss: 0.1095 - acc
uracy: 0.9682 - val_loss: 0.0946 - val_accuracy: 0.9743
Epoch 23/30
591/591 [=====] - 5s 8ms/step - loss: 0.1054 - acc
uracy: 0.9697 - val_loss: 0.0973 - val_accuracy: 0.9740
Epoch 24/30
591/591 [=====] - 5s 8ms/step - loss: 0.1027 - acc
uracy: 0.9700 - val_loss: 0.1036 - val_accuracy: 0.9741
Epoch 25/30
591/591 [=====] - 5s 8ms/step - loss: 0.1041 - acc
uracy: 0.9696 - val_loss: 0.1004 - val_accuracy: 0.9736
Out[16]: dict_keys(['loss', 'accuracy', 'val_loss', 'val_accuracy'])
```

In [17]:

```
import matplotlib.pyplot as plt

accuracy=history_out['accuracy']
accuracy_val=history_out['val_accuracy']


plt.clf()
plt.plot(accuracy,'b',label='training acc')
plt.plot(accuracy_val,'b', color='r', label='validation acc' )
plt.plot(test_acc, 'b', color='y', label='test acc')
plt.title('Training and validation accuracy')
plt.xlabel('epochs')
plt.ylabel('accuracy')
plt.legend()
plt.show()
```

```
/var/folders/9r/7f1bw3q15yz7435178lsgb5w0000gp/T/ipykernel_3430/2521518520.  
py:8: UserWarning: color is redundantly defined by the 'color' keyword argu  
ment and the fmt string "b" (-> color=(0.0, 0.0, 1.0, 1)). The keyword argu  
ment will take precedence.
```

```
plt.plot(accuracy_val, 'b', color='r', label='validation acc' )
```

```
/var/folders/9r/7f1bw3q15yz7435178lsgb5w0000gp/T/ipykernel_3430/2521518520.  
py:9: UserWarning: color is redundantly defined by the 'color' keyword argu  
ment and the fmt string "b" (-> color=(0.0, 0.0, 1.0, 1)). The keyword argu  
ment will take precedence.
```

```
plt.plot(test_acc, 'b', color='y', label='test acc')
```



It converges neatly without any vibration.