

# HW4 Pattern Recognition

2019150445 신백록

## 1. Download MNIST dataset

In [1]:

```
import numpy as np
from tensorflow.keras.datasets import mnist

(X,y),(X_test,y_test)=mnist.load_data()
print(X.shape)
print(X_test.shape)
print(y.shape)
print(y_test.shape)
```

```
Init Plugin
Init Graph Optimizer
Init Kernel
(60000, 28, 28)
(10000, 28, 28)
(60000,)
(10000,)
```

In [2]:

```
from sklearn.model_selection import train_test_split
X_train, X_val, y_train, y_val = train_test_split(X, y, test_size=0.1, stratify=y)
```

In [3]:

```
print(X_train.shape)
print(X_val.shape)
```

```
(54000, 28, 28)
(6000, 28, 28)
```

In [4]:

```
X_train=np.reshape(X_train,(-1,28,28,1))
X_test=np.reshape(X_test,(-1,28,28,1))
X_val=np.reshape(X_val,(-1,28,28,1))
```

In [5]:

```

from tensorflow.keras.utils import to_categorical

X_train=X_train/255.
X_val=X_val/255.
X_test=X_test/255.

y_train=to_categorical(y_train)
y_val=to_categorical(y_val)
y_test=to_categorical(y_test)
print(X_train.shape)
print(y_train.shape)

(54000, 28, 28, 1)
(54000, 10)

```

## 2. Build the CNN Model

In [6]:

```

from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten, Dense
from tensorflow.keras.models import Sequential
from tensorflow.keras.utils import plot_model

A_net=Sequential()
A_net.add(Conv2D(32,(3,3),activation='relu',input_shape=(28,28,1)))
A_net.add(MaxPooling2D((2,2), strides=2))
A_net.add(Conv2D(64,(3,3),activation='relu'))
A_net.add(MaxPooling2D((2,2), strides=2))
A_net.add(Conv2D(128,(3,3),activation='relu'))
A_net.add(MaxPooling2D((2,2), strides=2))
A_net.add(Flatten())
A_net.add(Dense(64,activation='relu'))
A_net.add(Dense(10,activation='softmax'))

A_net.summary()
plot_model(A_net,show_shapes=True)

```

Metal device set to: Apple M1  
Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
conv2d (Conv2D)	(None, 26, 26, 32)	320
-----		
max_pooling2d (MaxPooling2D)	(None, 13, 13, 32)	0
-----		
conv2d_1 (Conv2D)	(None, 11, 11, 64)	18496
-----		

```

max_pooling2d_1 (MaxPooling2 (None, 5, 5, 64))      0
conv2d_2 (Conv2D) (None, 3, 3, 128)      73856
max_pooling2d_2 (MaxPooling2 (None, 1, 1, 128))      0
flatten (Flatten) (None, 128)      0
dense (Dense) (None, 64)      8256
dense_1 (Dense) (None, 10)      650
=====
Total params: 101,578
Trainable params: 101,578
Non-trainable params: 0

```

```

2021-12-02 00:25:27.405725: I tensorflow/core/common_runtime/pluggable_device/pluggable_device_factory.cc:305] Could not identify NUMA node of platform GPU ID 0, defaulting to 0. Your kernel may not have been built with NUMA support.

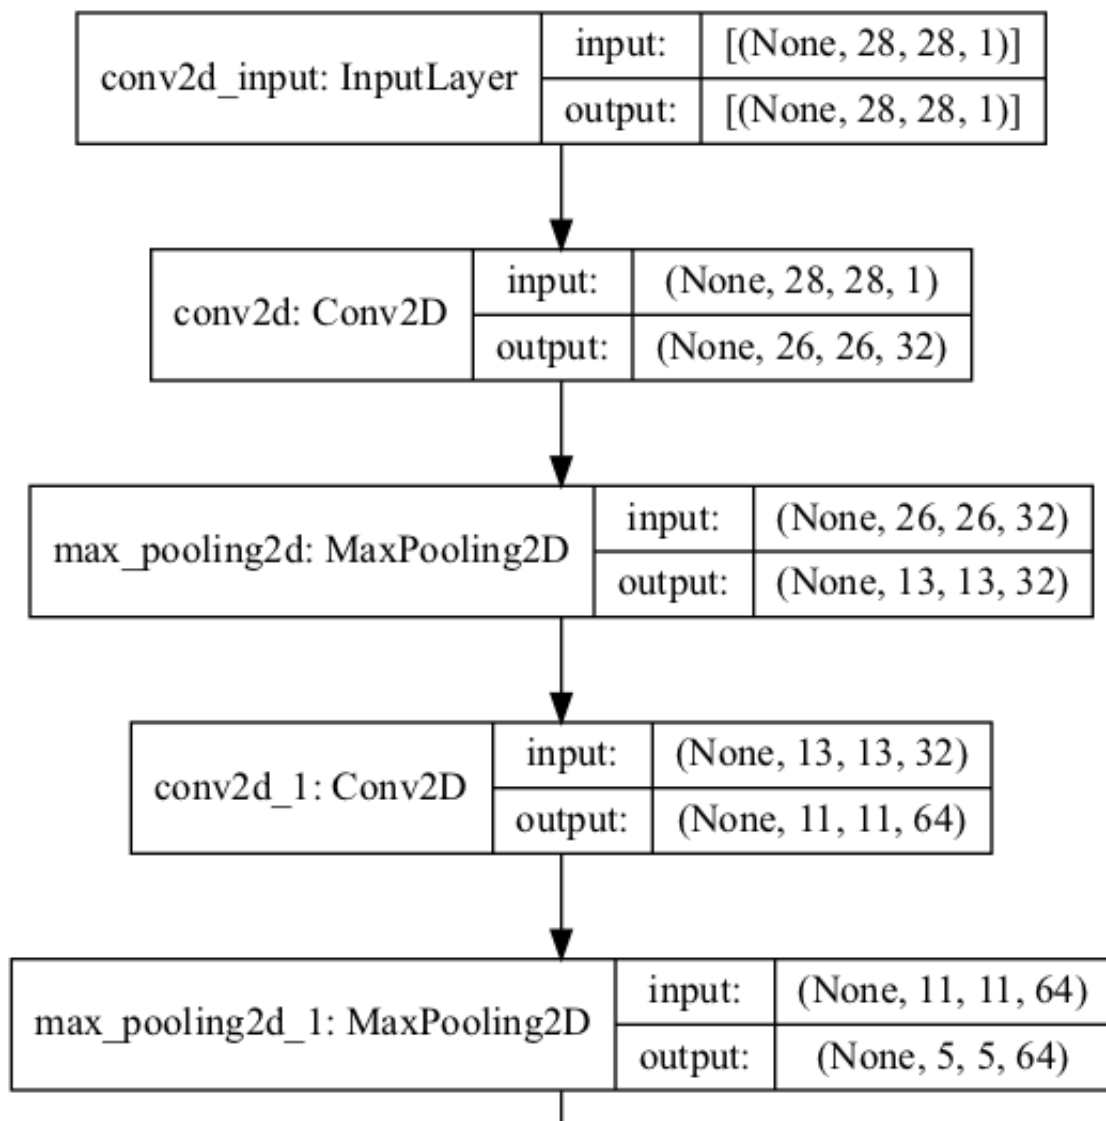
```

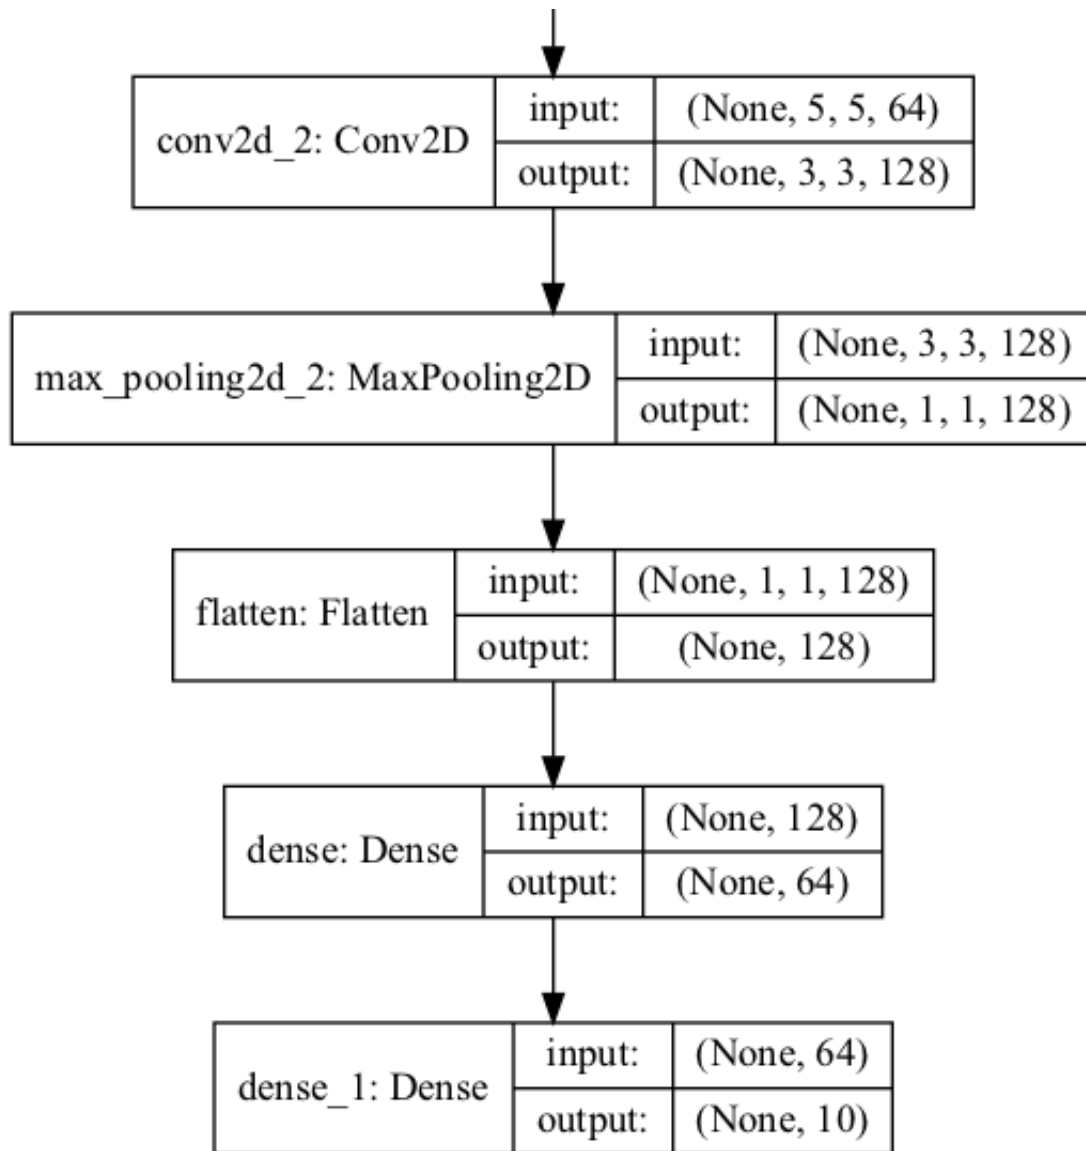
```

2021-12-02 00:25:27.405825: I tensorflow/core/common_runtime/pluggable_device/pluggable_device_factory.cc:271] Created TensorFlow device (/job:localhost/replica:0/task:0/device:GPU:0 with 0 MB memory) -> physical PluggableDevice (device: 0, name: METAL, pci bus id: <undefined>)

```

Out[6]:





In [7]:

```

B_net=Sequential()
B_net.add(Conv2D(16,(3,3),activation='relu',input_shape=(28,28,1),padding=
B_net.add(Conv2D(16,(3,3),activation='relu',padding='same'))
B_net.add(MaxPooling2D((2,2), strides=2))
B_net.add(Conv2D(32,(3,3),activation='relu'))
B_net.add(Conv2D(32,(3,3),activation='relu'))
B_net.add(MaxPooling2D((2,2), strides=2))
B_net.add(Flatten())
B_net.add(Dense(64,activation='relu'))
B_net.add(Dense(10,activation='softmax'))

B_net.summary()
plot_model(B_net, show_shapes=True)

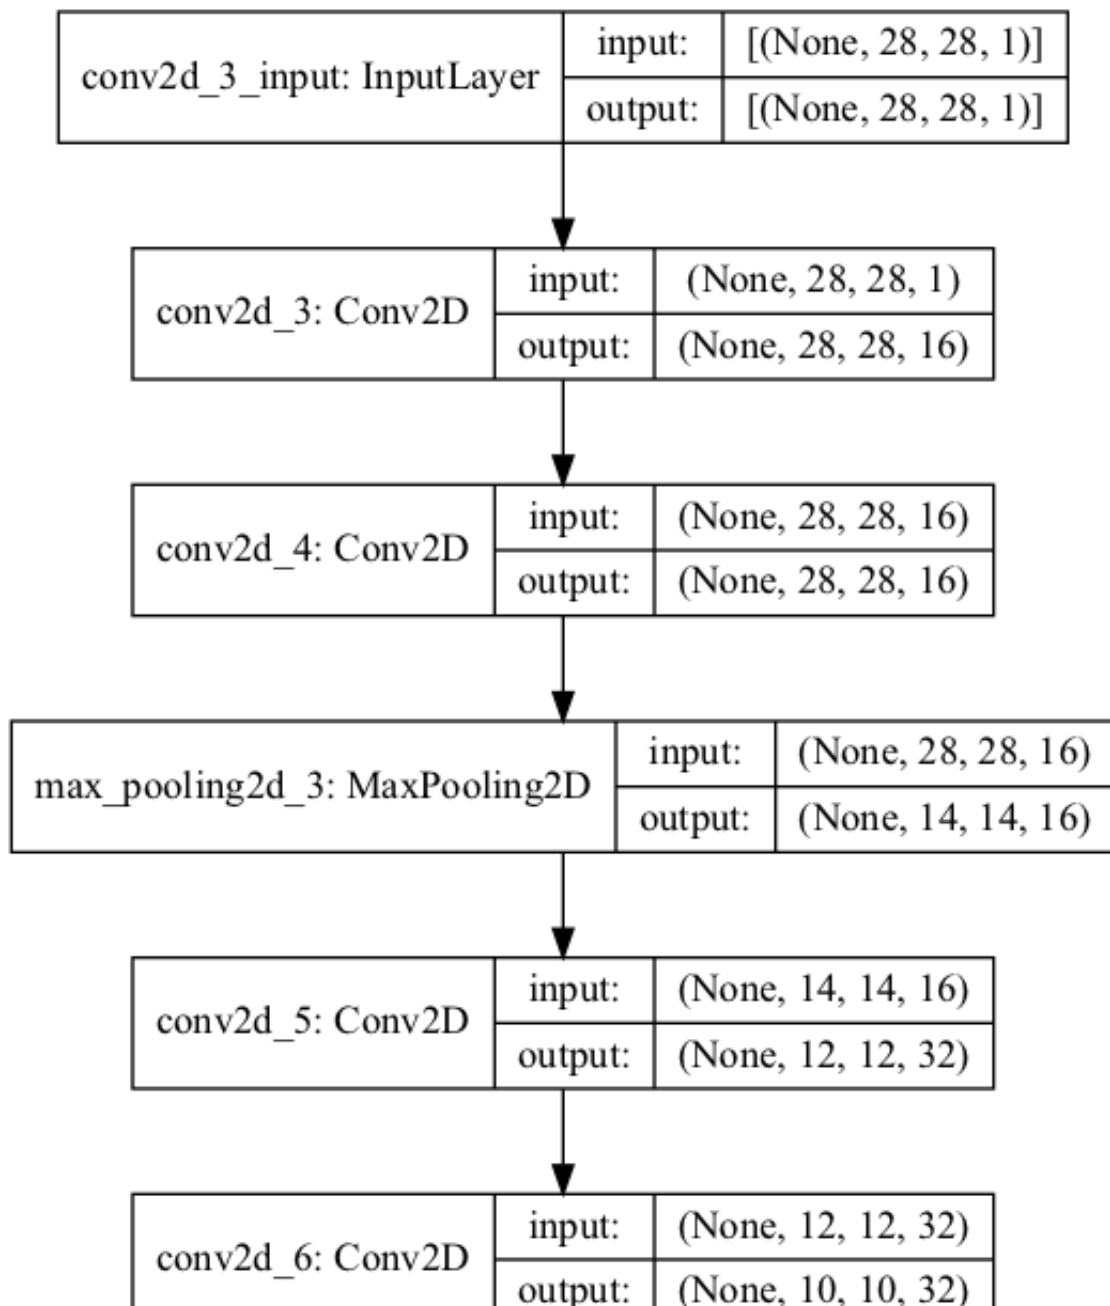
```

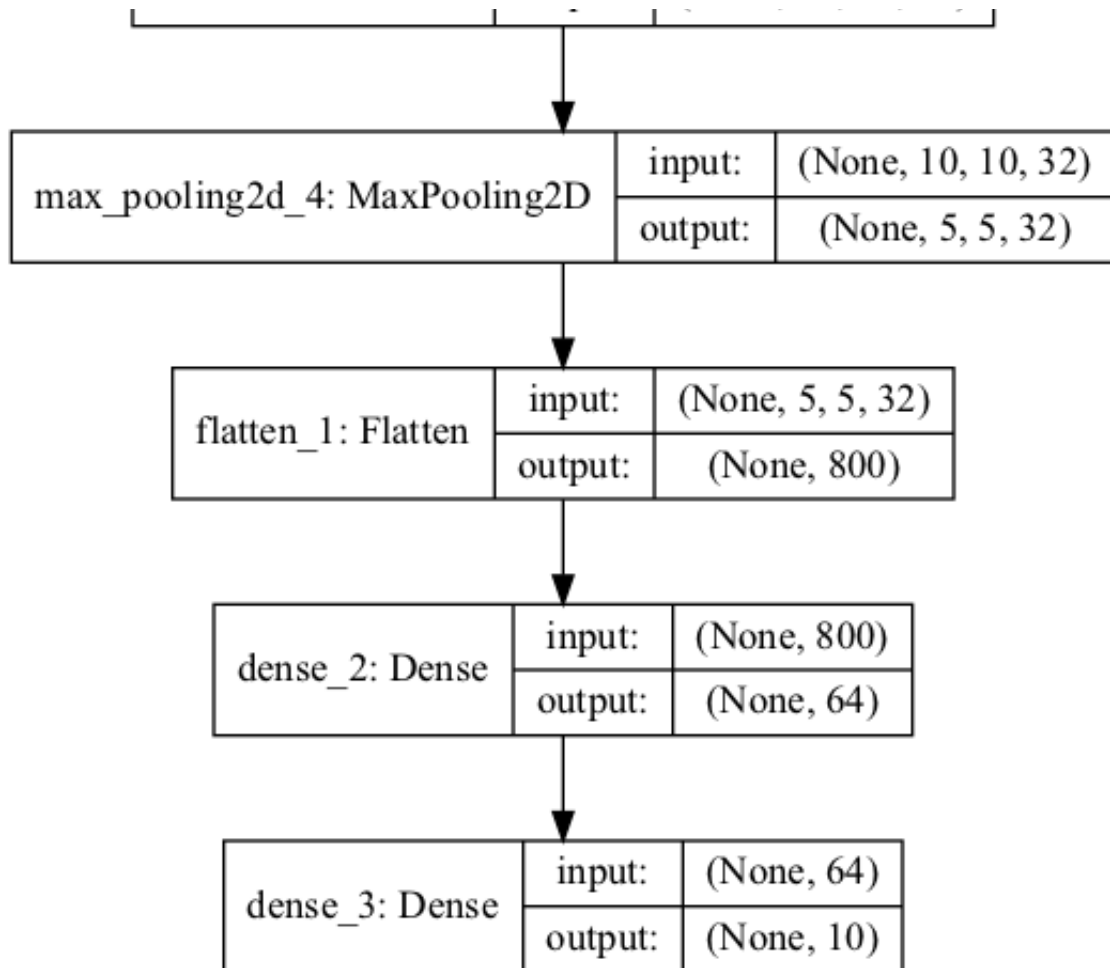
Model: "sequential\_1"

Layer (type)	Output Shape	Param #
=====		
conv2d_3 (Conv2D)	(None, 28, 28, 16)	160
=====		
conv2d_4 (Conv2D)	(None, 28, 28, 16)	2320

max_pooling2d_3 (MaxPooling2)	(None, 14, 14, 16)	0
conv2d_5 (Conv2D)	(None, 12, 12, 32)	4640
conv2d_6 (Conv2D)	(None, 10, 10, 32)	9248
max_pooling2d_4 (MaxPooling2)	(None, 5, 5, 32)	0
flatten_1 (Flatten)	(None, 800)	0
dense_2 (Dense)	(None, 64)	51264
dense_3 (Dense)	(None, 10)	650
=====		
Total params: 68,282		
Trainable params: 68,282		
Non-trainable params: 0		

Out[7]:





## 3. Training and Evaluation

### 3.a)

```
In [8]: A_net.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['accuracy'])
A_result=A_net.fit(X_train,y_train,batch_size=32,epochs=10,validation_data=(X_test,y_test))
```

```

2021-12-02 00:25:29.067786: I tensorflow/compiler/mlir/mlir_graph_optimization_pass.cc:176] None of the MLIR Optimization Passes are enabled (registered 2)
2021-12-02 00:25:29.071482: W tensorflow/core/platform/profile_utils/cpu_utils.cc:128] Failed to get CPU frequency: 0 Hz
2021-12-02 00:25:29.211361: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:112] Plugin optimizer for device_type GPU is enabled.
Epoch 1/10
1688/1688 [=====] - ETA: 0s - loss: 0.2151 - acc: 0.9339
2021-12-02 00:25:43.756813: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:112] Plugin optimizer for device_type GPU is enabled.
1688/1688 [=====] - 15s 9ms/step - loss: 0.2151 - acc: 0.9339 - val_loss: 0.0941 - val_acc: 0.9738
Epoch 2/10
1688/1688 [=====] - 15s 9ms/step - loss: 0.0747 - acc: 0.9763 - val_loss: 0.1026 - val_acc: 0.9693
Epoch 3/10
1688/1688 [=====] - 16s 9ms/step - loss: 0.0516 - acc: 0.9840 - val_loss: 0.0633 - val_acc: 0.9807
Epoch 4/10
1688/1688 [=====] - 16s 9ms/step - loss: 0.0408 - acc: 0.9875 - val_loss: 0.0595 - val_acc: 0.9828
Epoch 5/10
1688/1688 [=====] - 16s 9ms/step - loss: 0.0342 - acc: 0.9890 - val_loss: 0.0537 - val_acc: 0.9813
Epoch 6/10
1688/1688 [=====] - 16s 9ms/step - loss: 0.0259 - acc: 0.9916 - val_loss: 0.0406 - val_acc: 0.9890
Epoch 7/10
1688/1688 [=====] - 15s 9ms/step - loss: 0.0217 - acc: 0.9929 - val_loss: 0.0635 - val_acc: 0.9827
Epoch 8/10
1688/1688 [=====] - 15s 9ms/step - loss: 0.0182 - acc: 0.9941 - val_loss: 0.0447 - val_acc: 0.9877
Epoch 9/10
1688/1688 [=====] - 15s 9ms/step - loss: 0.0148 - acc: 0.9951 - val_loss: 0.0606 - val_acc: 0.9845
Epoch 10/10
1688/1688 [=====] - 16s 9ms/step - loss: 0.0135 - acc: 0.9956 - val_loss: 0.0579 - val_acc: 0.9860

```

In [9]:

```

A_val_loss, A_val_acc = A_net.evaluate(X_val,y_val)
print(A_val_acc)

```

```

188/188 [=====] - 1s 5ms/step - loss: 0.0579 - acc: 0.9860
0.9860000014305115

```

In [10]:

```

B_net.compile(loss='categorical_crossentropy',optimizer='adam', metrics=['accuracy'])
B_result=B_net.fit(X_train,y_train,batch_size=32,epochs=10,validation_data=(X_val,y_val))

```

```

Epoch 1/10
  6/1688 [.....] - ETA: 17s - loss: 2.2830 - acc: 0.1823
2021-12-02 00:28:13.572537: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:112] Plugin optimizer for device_type GPU is enabled.
1687/1688 [=====>.] - ETA: 0s - loss: 0.1448 - acc: 0.9538
2021-12-02 00:28:28.066847: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:112] Plugin optimizer for device_type GPU is enabled.
1688/1688 [=====] - 15s 9ms/step - loss: 0.1448 - acc: 0.9539 - val_loss: 0.0689 - val_acc: 0.9810
Epoch 2/10
1688/1688 [=====] - 16s 9ms/step - loss: 0.0461 - acc: 0.9856 - val_loss: 0.0504 - val_acc: 0.9855
Epoch 3/10
1688/1688 [=====] - 15s 9ms/step - loss: 0.0329 - acc: 0.9893 - val_loss: 0.0398 - val_acc: 0.9882
Epoch 4/10
1688/1688 [=====] - 16s 10ms/step - loss: 0.0248 - acc: 0.9921 - val_loss: 0.0441 - val_acc: 0.9883
Epoch 5/10
1688/1688 [=====] - 15s 9ms/step - loss: 0.0191 - acc: 0.9940 - val_loss: 0.0296 - val_acc: 0.9913
Epoch 6/10
1688/1688 [=====] - 16s 9ms/step - loss: 0.0172 - acc: 0.9945 - val_loss: 0.0350 - val_acc: 0.9908
Epoch 7/10
1688/1688 [=====] - 15s 9ms/step - loss: 0.0138 - acc: 0.9956 - val_loss: 0.0399 - val_acc: 0.9890
Epoch 8/10
1688/1688 [=====] - 16s 9ms/step - loss: 0.0122 - acc: 0.9960 - val_loss: 0.0374 - val_acc: 0.9897
Epoch 9/10
1688/1688 [=====] - 15s 9ms/step - loss: 0.0099 - acc: 0.9966 - val_loss: 0.0467 - val_acc: 0.9895
Epoch 10/10
1688/1688 [=====] - 15s 9ms/step - loss: 0.0089 - acc: 0.9971 - val_loss: 0.0381 - val_acc: 0.9903

```

In [11]:

```

B_val_loss, B_val_acc = B_net.evaluate(X_val,y_val)
print(B_val_acc)

```

```

188/188 [=====] - 1s 5ms/step - loss: 0.0381 - acc: 0.9903
0.9903333187103271

```

Since validation accuracy for A network is 0.986 & validation accuracy for B network is 0.9903, network B works slightly better than network A. But there could have some bias with the validation set so in order to precise the result, we have to consider doing cross validation or we have to analyze from test set. But in terms of the validation accuracy for this certain validation set, we can say network B is better.



### 3.b)

In [14]:

```
import matplotlib.pyplot as plt


accuracy=B_result.history['acc']
accuracy_val=B_result.history['val_acc']

plt.clf()
plt.plot(accuracy,'b',label='training acc')
plt.plot(accuracy_val,'b', color='r', label='validation acc' )

plt.title('Training and validation accuracy')
plt.xlabel('epochs')
plt.ylabel('accuracy')
plt.ylim(0.9,1)
plt.legend()
plt.show()
```

/var/folders/9r/7f1bw3q15yz7435178lsgb5w0000gp/T/ipykernel\_3857/2743705492.py:8: UserWarning: color is redundantly defined by the 'color' keyword argument and the fmt string "b" (-> color=(0.0, 0.0, 1.0, 1)). The keyword argument will take precedence.

```
plt.plot(accuracy_val,'b', color='r', label='validation acc' )
```



In first epoch, train accuracy is about 0.95 and validation accuracy is 0.981. When the epoch goes on, train accuracy and validation accuracy converges to 1 & about 0.99 respectively. We have to evaluate the model to test set to find whether the model is overfitted or not but in this graph with the validation accuracy, the model seems to be not overfitted or very slightly overfitted and it seems to converge well.

### 3.c)

In [15]:

```
B_net.evaluate(X_test, y_test)
```

```

313/313 [=====] - 2s 5ms/step - loss: 0.0306 - acc
: 0.9915
Out[15]: [0.030552800744771957, 0.9915000200271606]

```

Test accuracy is about 0.9915. Since the final train accuracy is 0.9971, we can say the model is very slightly overfitted or not overfitted to train set.

### 3.d)

First, add dropout function to the model between the layers to prevent overfitting. Second, add regularizer in Convolution layer and it will prevent overfitting problem. Also, Callback that stops learning when validation accuracy was not improved for several epochs would be helpful to prevent overfitting. Lastly, when we take a look at the B\_network, there are many parameters in Dense layer after Flatten function. So, we can use GlobalAveragePooling2D or GlobalMaxPooling2D instead of Flatten function and it makes the number of parameter much smaller than before.

For example in below, I add some Dropout layer between the Convolution layers and everything else is same with B\_net.

```

In [17]: from tensorflow.keras.layers import Dropout

B_net_adj=Sequential()
B_net_adj.add(Conv2D(16,(3,3),activation='relu',input_shape=(28,28,1),padding='same'))
B_net_adj.add(Conv2D(16,(3,3),activation='relu',padding='same'))
B_net_adj.add(MaxPooling2D((2,2), strides=2))
B_net_adj.add(Dropout(0.2))
B_net_adj.add(Conv2D(32,(3,3),activation='relu'))
B_net_adj.add(Dropout(0.2))
B_net_adj.add(Conv2D(32,(3,3),activation='relu'))
B_net_adj.add(MaxPooling2D((2,2), strides=2))
B_net_adj.add(Dropout(0.3))
B_net_adj.add(Flatten())
B_net_adj.add(Dense(64,activation='relu'))
B_net_adj.add(Dropout(0.4))
B_net_adj.add(Dense(10,activation='softmax'))

B_net_adj.summary()
plot_model(B_net_adj, show_shapes=True)

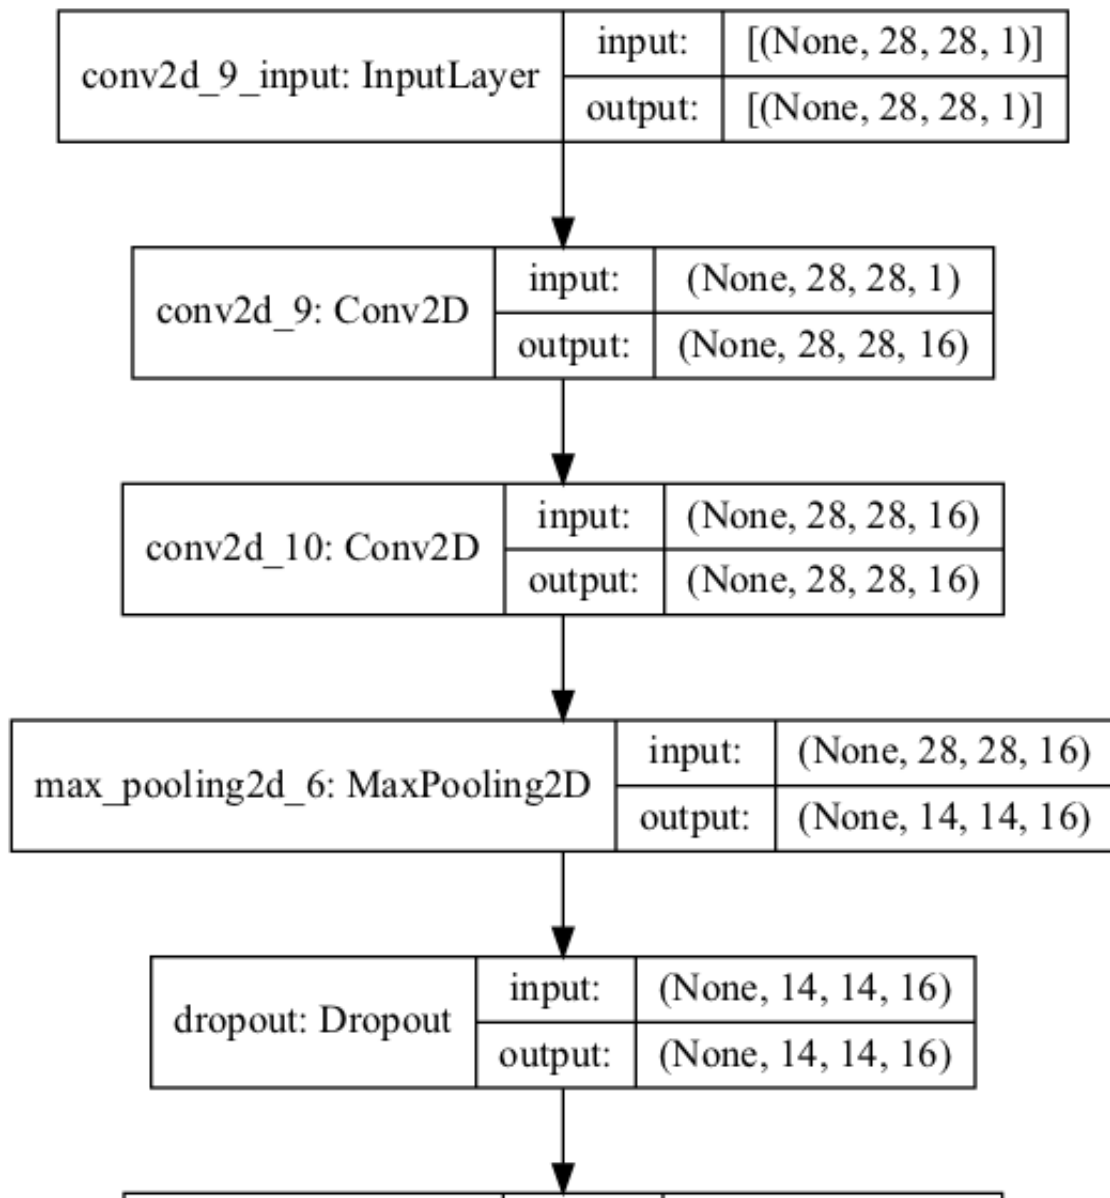
```

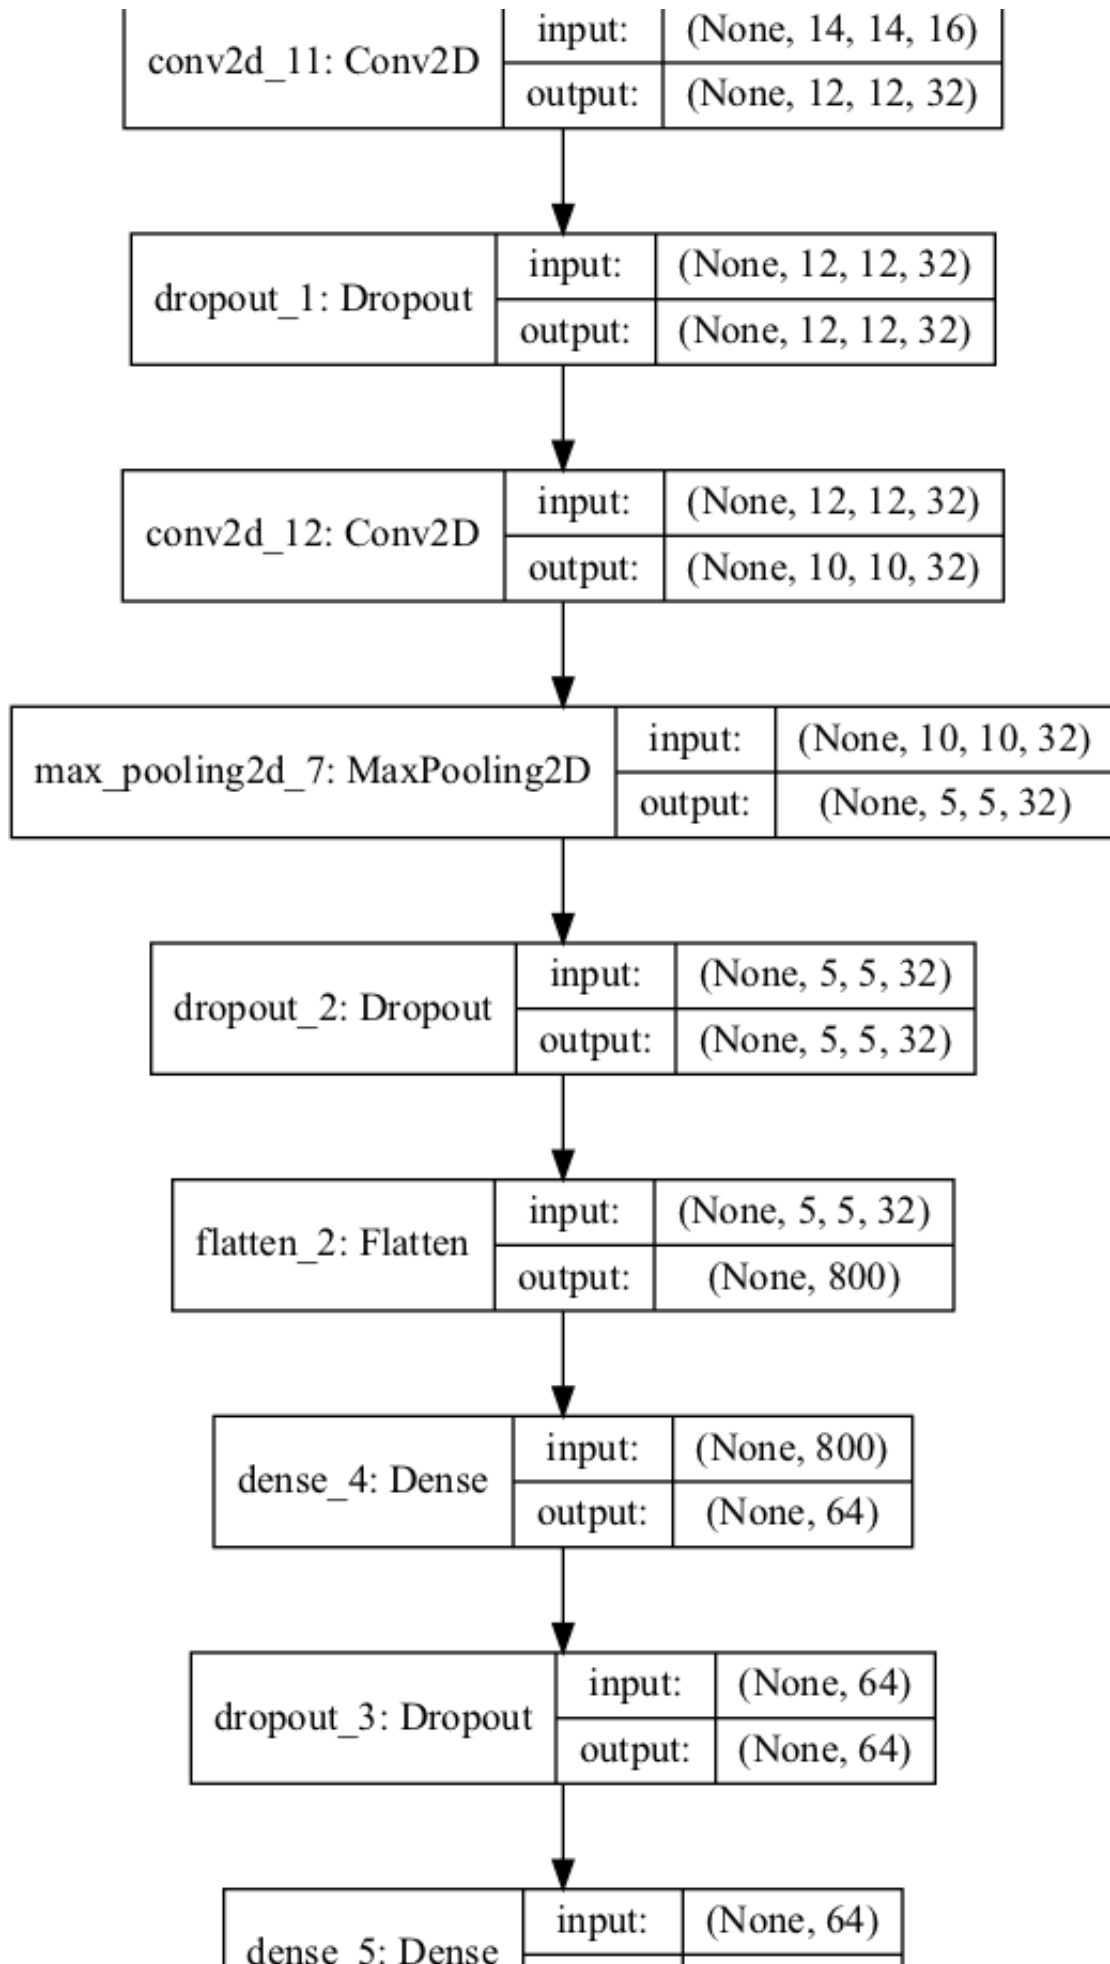
Model: "sequential\_3"

Layer (type)	Output Shape	Param #
conv2d_9 (Conv2D)	(None, 28, 28, 16)	160
conv2d_10 (Conv2D)	(None, 28, 28, 16)	2320
max_pooling2d_6 (MaxPooling2D)	(None, 14, 14, 16)	0
dropout (Dropout)	(None, 14, 14, 16)	0

conv2d_11 (Conv2D)	(None, 12, 12, 32)	4640
dropout_1 (Dropout)	(None, 12, 12, 32)	0
conv2d_12 (Conv2D)	(None, 10, 10, 32)	9248
max_pooling2d_7 (MaxPooling2D)	(None, 5, 5, 32)	0
dropout_2 (Dropout)	(None, 5, 5, 32)	0
flatten_2 (Flatten)	(None, 800)	0
dense_4 (Dense)	(None, 64)	51264
dropout_3 (Dropout)	(None, 64)	0
dense_5 (Dense)	(None, 10)	650
=====		
Total params: 68,282		
Trainable params: 68,282		
Non-trainable params: 0		

Out[17]:





-	output: (None, 10)
---	--------------------

In [18]:

```
B_net_adj.compile(loss='categorical_crossentropy',optimizer='adam', metrics=
B_result_adj=B_net_adj.fit(X_train,y_train,batch_size=32,epochs=10,validat
```

Epoch 1/10

```
2021-12-02 00:54:11.151232: I tensorflow/core/grappler/optimizers/custom_gr
aph_optimizer_registry.cc:112] Plugin optimizer for device_type GPU is enab
led.
```

```
1688/1688 [=====] - ETA: 0s - loss: 0.3370 - acc:
0.8929
```

```
2021-12-02 00:54:27.624037: I tensorflow/core/grappler/optimizers/custom_gr
aph_optimizer_registry.cc:112] Plugin optimizer for device_type GPU is enab
led.
```

```
1688/1688 [=====] - 18s 10ms/step - loss: 0.3370 -
acc: 0.8929 - val_loss: 0.0668 - val_acc: 0.9797
```

Epoch 2/10

```
1688/1688 [=====] - 17s 10ms/step - loss: 0.1221 -
acc: 0.9641 - val_loss: 0.0542 - val_acc: 0.9830
```

Epoch 3/10

```
1688/1688 [=====] - 17s 10ms/step - loss: 0.0962 -
acc: 0.9716 - val_loss: 0.0419 - val_acc: 0.9870
```

Epoch 4/10

```
1688/1688 [=====] - 17s 10ms/step - loss: 0.0791 -
acc: 0.9774 - val_loss: 0.0385 - val_acc: 0.9892
```

Epoch 5/10

```
1688/1688 [=====] - 18s 10ms/step - loss: 0.0717 -
acc: 0.9788 - val_loss: 0.0404 - val_acc: 0.9883
```

Epoch 6/10

```
1688/1688 [=====] - 18s 10ms/step - loss: 0.0656 -
acc: 0.9810 - val_loss: 0.0386 - val_acc: 0.9875
```

Epoch 7/10

```
1688/1688 [=====] - 17s 10ms/step - loss: 0.0598 -
acc: 0.9824 - val_loss: 0.0350 - val_acc: 0.9888
```

Epoch 8/10

```
1688/1688 [=====] - 17s 10ms/step - loss: 0.0578 -
acc: 0.9831 - val_loss: 0.0321 - val_acc: 0.9892
```

Epoch 9/10

```
1688/1688 [=====] - 17s 10ms/step - loss: 0.0541 -
acc: 0.9847 - val_loss: 0.0381 - val_acc: 0.9883
```

Epoch 10/10

```
1688/1688 [=====] - 17s 10ms/step - loss: 0.0534 -
acc: 0.9842 - val_loss: 0.0361 - val_acc: 0.9880
```

In [20]:

```
B_net_adj.evaluate(X_test, y_test)
```

```
313/313 [=====] - 2s 5ms/step - loss: 0.0219 - acc
: 0.9931
```

Out[20]: [0.021898901090025902, 0.9931000471115112]

The train accuracy in last epoch is about 0.984 but test accuracy is about 0.9931. It is better than B\_net in above! The Dropout function prevents overfitting and results the best output.

Lastly, I add some l2 regularizer in Dense layer after Flatten. Also I add Callback that stops learning if validation acc was not improved for 4 epochs.

In [33]:

```
import keras

B_net_adj2=Sequential()
B_net_adj2.add(Conv2D(16,(3,3),activation='relu',input_shape=(28,28,1),padding='same'))
B_net_adj2.add(Conv2D(16,(3,3),activation='relu',padding='same'))
B_net_adj2.add(MaxPooling2D((2,2), strides=2))
B_net_adj2.add(Dropout(0.2))
B_net_adj2.add(Conv2D(32,(3,3),activation='relu'))
B_net_adj2.add(Dropout(0.2))
B_net_adj2.add(Conv2D(32,(3,3),activation='relu'))
B_net_adj2.add(MaxPooling2D((2,2), strides=2))
B_net_adj2.add(Dropout(0.3))
B_net_adj2.add(Flatten())
B_net_adj2.add(Dense(64,activation='relu',kernel_regularizer=keras.regularizers.l2(0.01)))
B_net_adj2.add(Dropout(0.4))
B_net_adj2.add(Dense(10,activation='softmax'))

B_net_adj2.summary()
plot_model(B_net_adj2, show_shapes=True)
```

Model: "sequential\_8"

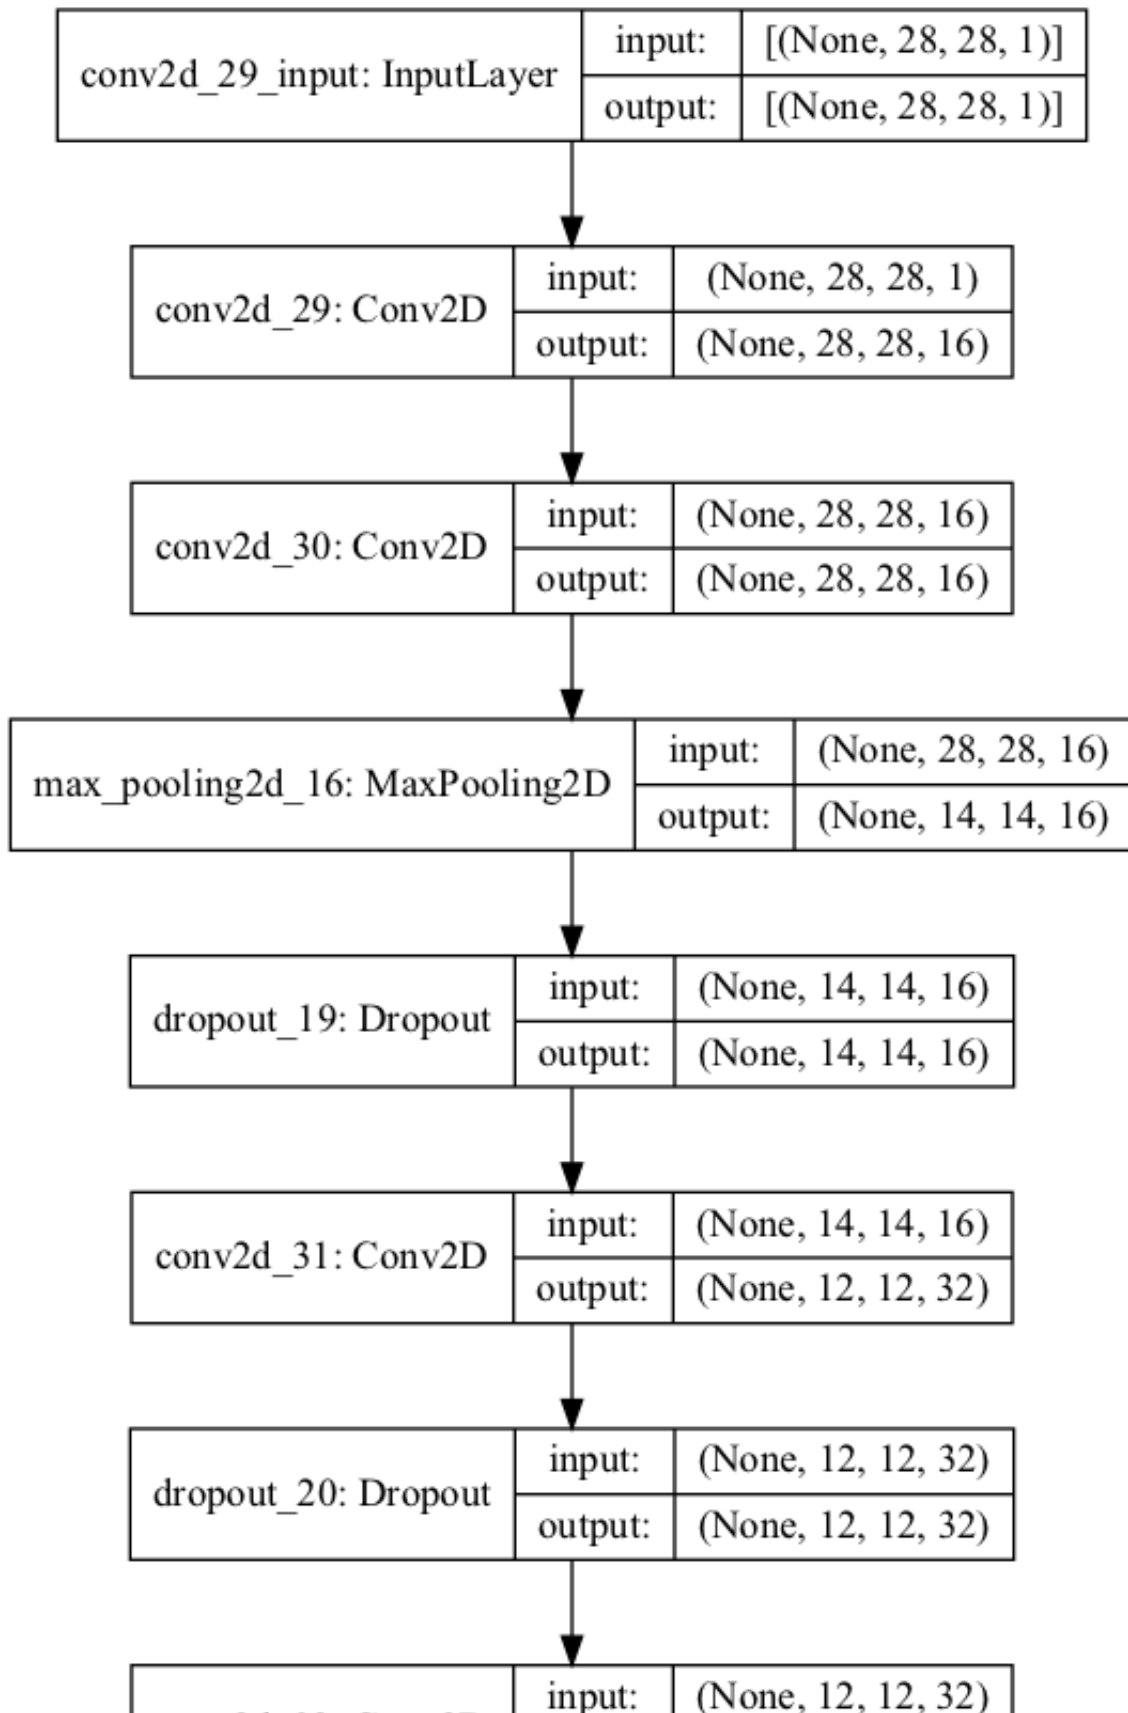
Layer (type)	Output Shape	Param #
=====		
conv2d_29 (Conv2D)	(None, 28, 28, 16)	160
conv2d_30 (Conv2D)	(None, 28, 28, 16)	2320
max_pooling2d_16 (MaxPooling)	(None, 14, 14, 16)	0
dropout_19 (Dropout)	(None, 14, 14, 16)	0
conv2d_31 (Conv2D)	(None, 12, 12, 32)	4640
dropout_20 (Dropout)	(None, 12, 12, 32)	0
conv2d_32 (Conv2D)	(None, 10, 10, 32)	9248
max_pooling2d_17 (MaxPooling)	(None, 5, 5, 32)	0
dropout_21 (Dropout)	(None, 5, 5, 32)	0
flatten_7 (Flatten)	(None, 800)	0
dense_12 (Dense)	(None, 64)	51264
dropout_22 (Dropout)	(None, 64)	0

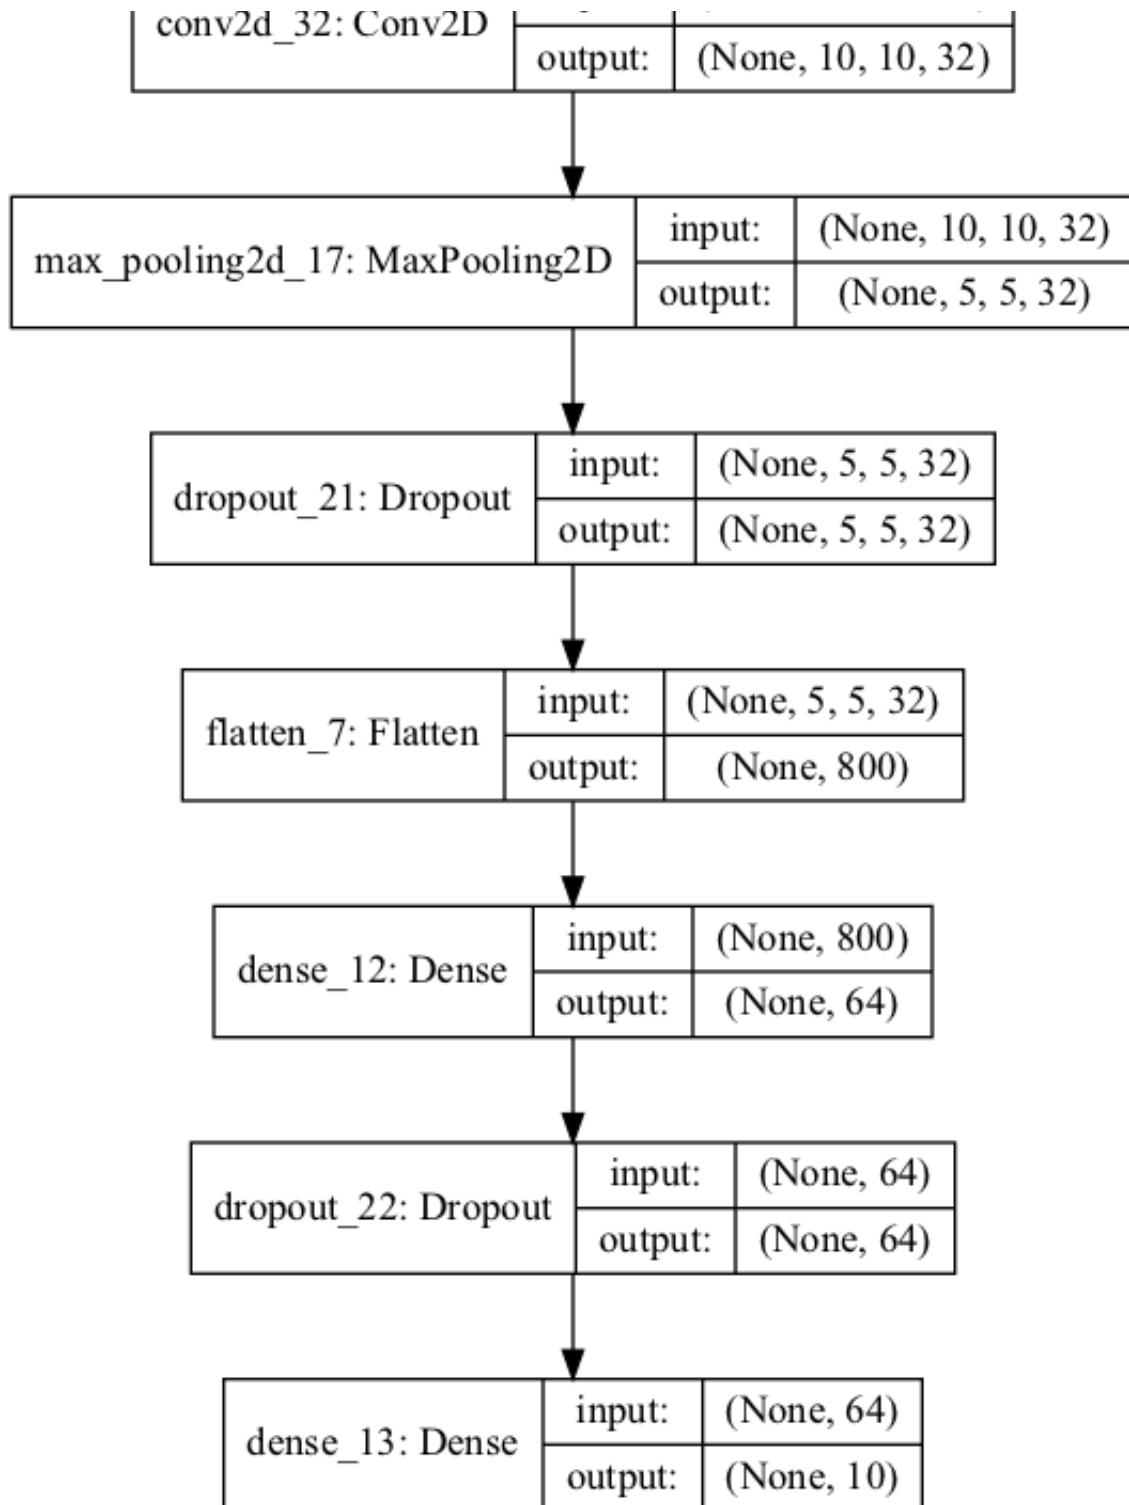
```

dense_13 (Dense)                (None, 10)                650
=====
Total params: 68,282
Trainable params: 68,282
Non-trainable params: 0

```

Out[33]:





```
In [36]: from tensorflow.keras.callbacks import EarlyStopping, ModelCheckpoint
callback_list=[EarlyStopping(monitor='val_acc',patience=4),ModelCheckpoint
```

```
In [37]: B_net_adj2.compile(loss='categorical_crossentropy',optimizer='adam', metric
B_result_adj2=B_net_adj2.fit(X_train,y_train,batch_size=32,epochs=20,callbacks
```

Epoch 1/20

3/1688 [.....] - ETA: 42s - loss: 0.3177 - acc: 0.9479



```
2021-12-02 01:17:05.890591: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:112] Plugin optimizer for device_type GPU is enabled.
1688/1688 [=====] - ETA: 0s - loss: 0.2544 - acc: 0.9587
2021-12-02 01:17:22.926795: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:112] Plugin optimizer for device_type GPU is enabled.
```

```

1688/1688 [=====] - 18s 11ms/step - loss: 0.2544 -
acc: 0.9587 - val_loss: 0.1750 - val_acc: 0.9837
Epoch 2/20
1688/1688 [=====] - 18s 11ms/step - loss: 0.2374 -
acc: 0.9642 - val_loss: 0.1692 - val_acc: 0.9832
Epoch 3/20
1688/1688 [=====] - 18s 11ms/step - loss: 0.2246 -
acc: 0.9663 - val_loss: 0.1648 - val_acc: 0.9845
Epoch 4/20
1688/1688 [=====] - 18s 11ms/step - loss: 0.2106 -
acc: 0.9703 - val_loss: 0.1588 - val_acc: 0.9848
Epoch 5/20
1688/1688 [=====] - 18s 11ms/step - loss: 0.2004 -
acc: 0.9718 - val_loss: 0.1566 - val_acc: 0.9850
Epoch 6/20
1688/1688 [=====] - 18s 11ms/step - loss: 0.1967 -
acc: 0.9726 - val_loss: 0.1439 - val_acc: 0.9878
Epoch 7/20
1688/1688 [=====] - 18s 11ms/step - loss: 0.1920 -
acc: 0.9741 - val_loss: 0.1450 - val_acc: 0.9878
Epoch 8/20
1688/1688 [=====] - 18s 11ms/step - loss: 0.1882 -
acc: 0.9750 - val_loss: 0.1448 - val_acc: 0.9860
Epoch 9/20
1688/1688 [=====] - 18s 11ms/step - loss: 0.1862 -
acc: 0.9748 - val_loss: 0.1433 - val_acc: 0.9872
Epoch 10/20
1688/1688 [=====] - 18s 11ms/step - loss: 0.1797 -
acc: 0.9761 - val_loss: 0.1312 - val_acc: 0.9888
Epoch 11/20
1688/1688 [=====] - 18s 11ms/step - loss: 0.1786 -
acc: 0.9768 - val_loss: 0.1377 - val_acc: 0.9887
Epoch 12/20
1688/1688 [=====] - 18s 11ms/step - loss: 0.1771 -
acc: 0.9771 - val_loss: 0.1383 - val_acc: 0.9878
Epoch 13/20
1688/1688 [=====] - 18s 11ms/step - loss: 0.1729 -
acc: 0.9775 - val_loss: 0.1237 - val_acc: 0.9905
Epoch 14/20
1688/1688 [=====] - 18s 11ms/step - loss: 0.1730 -
acc: 0.9786 - val_loss: 0.1326 - val_acc: 0.9885
Epoch 15/20
1688/1688 [=====] - 18s 11ms/step - loss: 0.1712 -
acc: 0.9777 - val_loss: 0.1333 - val_acc: 0.9885
Epoch 16/20
1688/1688 [=====] - 18s 11ms/step - loss: 0.1694 -
acc: 0.9778 - val_loss: 0.1222 - val_acc: 0.9897
Epoch 17/20
1688/1688 [=====] - 18s 11ms/step - loss: 0.1687 -
acc: 0.9787 - val_loss: 0.1228 - val_acc: 0.9898

```

```
In [38]: B_net_adj2.evaluate(X_test, y_test)
```

```
313/313 [=====] - 2s 6ms/step - loss: 0.1113 - acc  
: 0.9934  
Out[38]: [0.11134836822748184, 0.9934000372886658]
```

The train accuracy for last epoch is 0.979 but test accuracy is 0.9934. It is slightly better than above model!